

RUCHIR ATTRI

Toronto, ON | 236 862-7159 | ruchir.attri99@gmail.com | [LinkedIn](#) | [Github](#)

Professional Summary

- AI & Automation Specialist:** 3+ years of experience engineering scalable AI solutions and NLP pipelines using Python, spaCy, and SQL for Fortune 500 clients.
- LLM & RAG Developer:** Proven ability to architect autonomous agents and RAG systems using LangChain, FastMCP, and Vector DBs, deploying on platforms like GCP Vertex AI.
- Data Engineering Expert:** Skilled in processing large-scale datasets and building automated extraction workflows for unstructured document analysis.

Technical Skills

| | |
|------------------------|--|
| Languages & Frameworks | Python, SQL, FastAPI, LangChain, spaCy, Pydantic, Node.js (Basic) |
| AI & ML | LLMs (Llama, GPT), RAG, NLP, OCR, XGBoost, Computer Vision, Prompt Engineering |
| Cloud & DevOps | GCP (Vertex AI), AWS, Azure, Docker, Git, CI/CD |
| Data & Databases | Redis, Vector Databases, SAS Enterprise Miner, PostgreSQL |

Professional Experience

| | |
|--|---|
| Canadian Imperial Bank of Commerce Data Scientist | September 2024 – December 2024 Toronto, ON |
| <ul style="list-style-type: none">Engineered high-dimensional training datasets with 250+ features by joining 14M+ client records and marketing signals using SAS SQL, enabling robust credit risk modeling.Developed explainable decision tree and ensemble classifiers in SAS Enterprise Miner to segment clients for credit products | |
| Accenture Machine Learning Engineer | July 2021 - July 2023 Bengaluru, India |
| <ul style="list-style-type: none">Architected a PII detection pipeline using Python and spaCy NLP, increasing compliance reporting accuracy by 60% and automating sensitive data protection for Fortune 500 clients.Built 10+ automated data connectors to parse and extract structured insights from unstructured documents across AWS, Azure, SharePoint, and Jira, streamlining data ingestion.Optimized an XGBoost classification framework by aggregating confidence scores from multiple models, significantly enhancing prediction reliability for critical business processes.Led cross-functional engineering collaboration, managing weekly deliverables and resolving technical blockers to ensure on-time product deployment.Diagnosed and resolved complex client-reported issues for a Data Governance SaaS product, tailoring solutions to meet specific Fortune 500 operational needs. | |

Projects

| | |
|---|---------------|
| CKAN MCP Server | November 2025 |
| <ul style="list-style-type: none">Championed the adoption of Agentic AI within the public sector by successfully pitching and demonstrating the MCP server to the Toronto Open Data Team, establishing a proof-of-concept for future AI-driven civic technology.Architected a Model Context Protocol (MCP) server in Python to bridge Large Language Models (LLMs) with CKAN open data portals, enabling autonomous agents to perform complex data discovery and retrieval tasks without manual intervention.Engineered 13+ asynchronous tools and context-aware prompts using FastMCP and Pydantic, facilitating advanced agentic workflows such as schema extraction, SQL-like datastore querying, and intelligent resource previewing. | |
| GPT From Scratch | August 2025 |
| <ul style="list-style-type: none">Constructed a decoder-only transformer model from first principles in PyTorch, demonstrating deep understanding of attention mechanisms and LLM architecture. | |
| WikiHouw RAG Application | May 2025 |
| <ul style="list-style-type: none">Developed a Retrieval Augmented Generation (RAG) application using LangChain and Llama 3.2, enabling context-aware query answering for complex documentation.Implemented Redis for LLM memory and caching, optimizing retrieval latency and improving response quality for user queries. | |

Google Hackathon

February 2024

- Deployed sales forecasting models using GCP Vertex AI (AutoML), identifying 50% excess inventory risk and enabling proactive supply-chain automation.
- Integrated marketing personalization workflows via Vertex AI, utilizing prompt engineering to automate and tailor customer outreach strategies.
- Designed a comprehensive Business Pulse Dashboard on Looker Studio for real-time monitoring of sales and product performance metrics.

Education

Master of Business Analytics

August 2023 - September 2024

UBC Sauder School of Business

Vancouver, BC

Key Courses: Marketing Analytics, Customer Analytics, Machine Learning, Analytics Leadership, Financial Analytics, Optimization Techniques, Database Management Systems, SQL, Experimentation and Causal Inferences, Data Visualization, Statistical Analysis

Bachelor of Technology in **Electronics & Telecommunication Engineering**

July 2017 - July 2021

KJ Somaiya College of Engineering, University of Mumbai

Mumbai, India

Key Courses: Neural Networks and Fuzzy Logic, Image Processing, Data Structures and Analysis of Algorithms